

# **Learning SQL the Badass Way**

Edgar J. Treischl

# Table of contents

<b>Preface</b>	<b>4</b>
<b>1 Introduction</b>	<b>7</b>
1.1 Select . . . . .	8
1.2 Where . . . . .	9
1.3 Count . . . . .	10
1.4 DISTINCT . . . . .	11
1.5 Insert Values . . . . .	11
1.6 Updates . . . . .	12
1.7 Delete . . . . .	13
<b>2 Data management</b>	<b>14</b>
2.1 Types of SQL Statements . . . . .	15
2.2 CREATE . . . . .	15
2.3 ALTER TABLE . . . . .	16
2.4 Truncate . . . . .	19
<b>3 Calculations</b>	<b>20</b>
3.1 String values, ranges and set of values . . . . .	20
3.1.1 Sorting Result Sets . . . . .	21
3.1.2 Grouping result sets . . . . .	23
3.2 Built-in Database Functions . . . . .	25
3.2.1 Date and Time Built-in Functions . . . . .	26
3.2.2 Sub-Queries and Nested Selects . . . . .	27
3.2.3 Working with Multiple Tables . . . . .	29
3.3 Summary . . . . .	31
<b>4 Cheats</b>	<b>32</b>
4.1 Introduction . . . . .	32
4.1.1 Select . . . . .	32
4.1.2 Where . . . . .	32
4.1.3 Count . . . . .	33
4.1.4 DISTINCT . . . . .	33
4.1.5 Insert Values . . . . .	33
4.1.6 Updates . . . . .	34
4.1.7 Delete . . . . .	34

4.2	Data management . . . . .	34
4.2.1	Create table . . . . .	34
4.2.2	Alter table . . . . .	35
4.2.3	Truncate . . . . .	37
4.3	Calculations . . . . .	37
4.3.1	Like . . . . .	37
4.3.2	Between . . . . .	38
4.3.3	In . . . . .	38
4.3.4	Order by . . . . .	39
4.3.5	Distinct . . . . .	39
4.3.6	Group by . . . . .	39
4.3.7	Average . . . . .	40
4.3.8	Having . . . . .	40
4.3.9	Further functions . . . . .	40
4.3.10	Date and Time Built-in Functions . . . . .	42
4.3.11	Working with Multiple Tables . . . . .	42

<b>References</b>	<b>45</b>
-------------------	-----------

# Preface

*Learning SQL the badass way* is an unpublished manuscript which I started to summarize the content of an SQL course. Please keep that in mind when reading this book.

Learning SQL the *badass way* shows you the basic commands to manage data with SQL. I wrote this manual because learning new things takes me a lot of time and is often designed without taking our prior knowledge about data science into account. I started to learn SQL via an online course, but instead of digging myself into the applied side of handling data with SQL, I got a lot of information about SQL in general terms, data bases, and other related topics. However, I was able to read and to write SQL even when I had no prior experience working with SQL because I have a background in R (and Python) and SQL is not rocket science. Thus, this manual is the result from taking this course, which helps me when needed and it may help you to quickly learn SQL if you have a similar background,

The manual summarizes the content of the [IBM SQL for Data Science Course \(from EDx\)](#) based on own examples and data which I attained in 2022. It's a quick and dirty introduction for R Users: I assume that you are familiar with R and I will skip many typical steps of a real SQL introduction; certainly I will not explain why we need to wrangle data in the first place. This manual only summarizes main idiosyncrasies of SQL, not concepts that you probably know from any other programming experience. Please, close this manual if you want to learn SQL in a proper way. Go and attend a real course and where are many books and other resources available. However, if you are bored from long introduction what data is, why we need to learn how to wrangle data or other common aspects that come along the data science journey, feel free to join me.

Learning SQL is beneficial but maybe we need more motivational input. ChatGTP gives us the following reasons why we should learn SQL even if you are fluent in a language such as R:

- *Efficient data management*: SQL is designed to work with relational databases, which are ideal for managing large amounts of structured data. By learning SQL, R users can efficiently query, retrieve and manage data from databases using SQL commands, making their data analysis tasks more efficient.
- *Collaborative work*: SQL is a common language used by data analysts and data engineers, making it easier to collaborate and share data between different teams. By learning SQL, R users can communicate more effectively with other data professionals and work collaboratively on projects.

- *Integration with R:* R users often work with data that is stored in databases, and SQL provides a way to query and retrieve data from these databases directly into R. This integration allows R users to take advantage of the power of SQL for data management, while still working with their preferred R environment.
- *Advanced data manipulation:* SQL provides powerful features for data manipulation, including filtering, sorting, aggregating and joining data from multiple tables. By learning SQL, R users can take advantage of these advanced features to manipulate their data in more sophisticated ways.
- *Job market demand:* SQL is a widely used skill in the data analytics job market. By learning SQL, R users can broaden their skillset and increase their job market competitiveness.

ChatGTP also helps us to get an overview about the most important R packages that are commonly used for working with SQL, including:

- **dplyr:** A powerful package for data manipulation that can connect to various SQL databases and perform operations such as filtering, grouping, and joining (Wickham et al. 2023).
- **DBI:** An R package that provides a common interface for connecting to various SQL databases (, Wickham, and Müller 2022).
- **RMySQL and RSQLite:** Packages that provide an interface for connecting to MySQL and SQLite databases, respectively (Ooms et al. 2022; Müller et al. 2023).
- **RJDBC:** A package that provides a JDBC interface for connecting to various databases, including Oracle, Microsoft SQL Server, and PostgreSQL (Urbanek 2022).
- **RODBC:** A package that provides an ODBC interface for connecting to various databases, including Microsoft SQL Server and PostgreSQL (Ripley and Lapsley 2022).
- **sqldf:** A package that allows you to run SQL queries on data frames in R (Grothendieck 2017).

Overall, the choice of which package to use will depend on the specific database you're working with and your preferred interface. However, these packages should provide a good starting point for working with SQL in R and I will outline more about each package soon. Before we start can start with SQL, let us set the scope of this manual:

- Chapter 1 introduces the basics SQL commands and shows you how to run SQL queries from the R console.
- Chapter 2 elaborates on SQL knowledge in terms of data management and introduces, for example, two different types of SQL statements.

- Chapter 3 shows how basic calculations are done in SQL: We can sort and group data or work with multiple tables.
- Chapter 4 is a cheat sheet. It repeats all introduced SQL commands based on a simple code snippets. By doing so, we can copy and paste the code if needed.

# 1 Introduction

The biggest mistake trying to learn SQL first. If you ever opened a book that introduces SQL (or any other programming language) you will find several chapters that outline what SQL (structured query language) is. An explanation from the Wikipedia page should be sufficient: “SQL a domain-specific language used in programming and designed for managing data held in a relational database management system (RDBMS), or for stream processing in a relational data stream management system (RDSMS)”. Thus, in a nutshell we can remember that SQL is a language to manage data (tables) in databases.

The next chapters of a typical introduction shows you how to setup the software and the database. I'll skip this step because setting up a database on your own computer does not make much sense and there are smarter ways to learn the first steps if you are a R user. For example, this website includes SQL code and output but this does not imply that I have connected to a database. Instead we can use R packages such as DBI to setup the connection and include your SQL code directly in your RMarkdown document or an R script. There is no need for a database in the beginning, we can make use of the *local memory* to simulate a database, save a table (data frame) and run SQL commands directly as a code chunk from the R script. Consider the next console, it shows SQL code to select a variable from a data set as an example.

```
SELECT VAR FROM data;
```

To run this code from R without a database, we need to establish a local connection (`con`), write the data into the local data base with `dbWriteTable()`. The next console shows how it works.

```
library(DBI)
library(RSQLite)

# Create in-memory RSQLite database
con <- dbConnect(drv = RSQLite::SQLite(),
                 dbname = ":memory:")

#Write a table into the data base
dbWriteTable(conn = con,
             name = "mtcars",
```

```
value = mtcars)
```

I hope you are familiar with the `mtcars` and the `iris` data if you want to reproduce these first SQL steps. Both are implemented in R and in this chapter we explore first SQL commands with this data. To this end, you need to insert a table into the data base. As the last console for the `mtcars` data highlighted. The `dbListTables()` functions lists all loaded tables of the used data base.

```
dbListTables(con)
```

```
[1] "df"      "iris"    "mtcars"
```

After we have made a connection and the data is available, the `sqldf` package let us run SQL code from R.

```
#Run SQL code from R
library(sqldf)
sqldf('SELECT * FROM mtcars LIMIT 5;')
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
2	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
3	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
4	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
5	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2

Thus, such R packages increases your learning curve substantially, since we can establish a data base and try some SQL code without any other equipment than your local machine. In Chapter XX we will establish a connection to a SQL data base, but for the moment there is no need to. The remaining part of this chapter shows first SQL commands.

## 1.1 Select

- Use select to retrieve a table or a column from a table
- You can select a single column from a table
- Or select the entire table (data frame) with the wildcard `*`



```
SELECT * FROM mtcars;
```

Table 1.1: Displaying records 1 - 10

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4

- Limit the output by providing the number of lines

```
SELECT mpg, disp FROM mtcars LIMIT 5;
```

Table 1.2: 5 records

mpg	disp
21.0	160
21.0	160
22.8	108
21.4	258
18.7	360

- You can also insert a starting point that skips some observations. For instance, `OFFSET 10` will skip the first ten table entries
- Use quotations marks if the column contains special characters (like `'Petal.Width'` from iris data)

## 1.2 Where

- Define what you want to select with the `Where` option (SQL folks say clause)
- For instance, the variable `am` is a binary indicator (0/1) and you can use *where* to select data only if `am = 0`

```
SELECT * FROM mtcars WHERE am = 0 LIMIT 5;
```

Table 1.3: 5 records

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2

- Remember to use quotation marks if you try to use *where* with non-numerical values from: e.g. != 'label'

```
SELECT * FROM iris WHERE Species = "virginica" LIMIT 5;
```

Table 1.4: 5 records

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
6.3	3.3	6.0	2.5	virginica
5.8	2.7	5.1	1.9	virginica
7.1	3.0	5.9	2.1	virginica
6.3	2.9	5.6	1.8	virginica
6.5	3.0	5.8	2.2	virginica

## 1.3 Count

- Count counts cases!

```
SELECT COUNT(*) FROM mtcars;
```

Table 1.5: 1 records

COUNT(*)
32

- We can count subgroups with the **WHERE** clause:

```
SELECT COUNT(am) FROM mtcars WHERE am != 0;
```

Table 1.6: 1 records

COUNT(am)
13

## 1.4 DISTINCT

- `Distinct` can be used to find distinct values. For instance, there are three different species in the iris data:

```
SELECT DISTINCT Species FROM iris
```

Table 1.7: 3 records

Species
setosa
versicolor
virginica

- As in other programming languages, we can combine several commands. For instance, we can `COUNT` how many `Distinct` species the iris data has:

```
SELECT COUNT (DISTINCT Species) FROM iris
```

Table 1.8: 1 records

COUNT (DISTINCT Species)
3

## 1.5 Insert Values

- Next, I use a small data set (`df`) to illustrate how to *insert values*, *make updates*, and *delete cases*
- My toy data set `df` has two observations with three variables: `x,y,z`

- Never mind if you do not know what a `tribble` is, it is just a command to create data

```
library(tidyverse)

df <- tribble(
  ~x, ~y, ~z,
  "a", 2, 3.6,
  "b", 1, 8.5
)
df
```

```
# A tibble: 2 x 3
  x     y     z
<chr> <dbl> <dbl>
1 a     2   3.6
2 b     1   8.5
```

- Now, we can insert new values into `df` by providing a list of the columns you want to fill in with values for each column:

```
INSERT INTO df (x, y, z) VALUES('c', 3, 1);
```

Let's see whether it worked:

```
SELECT * FROM df;
```

Table 1.9: 3 records

x	y	z
a	2	3.6
b	1	8.5
c	3	1.0

## 1.6 Updates

- Make updates for single (or multiple) values
- For instance, we can update the variable `z` and set `z = 77` for a certain level of another variable:

```
UPDATE df SET z = 77 WHERE x = 'b';
```

- Take care, without the **WHERE** clause all observation would get the new value!

```
SELECT * FROM df;
```

Table 1.10: 3 records

x	y	z
a	2	3.6
b	1	77.0
c	3	1.0

## 1.7 Delete

- We can drop or delete observations, but of course we should take care since we probably do not want to delete the entire table, just for some implausible values
- For this reason we use the **WHERE** clause again, for instance, to get rid of second row of the toy data set:

```
DELETE from df WHERE x = 'b';
```

```
SELECT * FROM df;
```

Table 1.11: 2 records

x	y	z
a	2	3.6
c	3	1.0

In summary, we have to select a table from the database, specify conditions with the where clause and we can use count to get a first impression of the data. We saw how to insert values, a really vague concept if you are used to work with data, but from a SQL perspective you give the database a update or imagine that a stream of new data needs an update. If SQL seems like something you have to get used to it, don't be afraid, me too.

## 2 Data management

This is the second blog post from “Learn SQL the badass way”. I outlined the scope of the blog in the first blog entry where I explained why I set the scope to R users and people with some programming languages. You may not find what you are looking for if you are not familiar with R or if you do not have any other experience to work with data, because I focus on the data management part in this blog.

Before we get in touch with new SQL commands, we have to learn some SQL vocabulary. As in other languages, we have to learn some basic vocabulary to advance our SQL skills. Let’s say we have a small collection of tables about books. In the **author** table we store information about the book **author**(e.g. first name, second name, etc); in the **book** table contains information about each book (e.g. genre etc); and the **sales** table summarizes the sales data for each book. How do we manage all of these tables and the dependencies in SQL?

Sometimes you encounter a diagram to display how each table is related to each other and we can think of all the independent tables as a collection of tables and we have to figure out how they are related. A entity relationship diagram (ERD) will help you to see the relation of each table, it displays the collection of entities and attributes.

In the SQL world, *entities* are independent objects. For instance, the book table is a independent object because it exists next to other entities of our collection. Entities have *attributes* or properties. For instance, the book table contains title, subtitle, book id and more attributes. Thus, entities refer to tables (or in my world data frame) of our collection and attributes refer to columns, or I would say variables, in the table.

Furthermore, we can differentiate between primary and foreign keys in the tables:

- *Primary key*: Is a unique indicator that helps us to match tables (e.g. a unique author ID)
- *Foreign keys*: Is a primary key that is defined in other table to create a link between tables (Book ID in Book table and the sales table)

We may also differentiate what information a attribute stores. Some common data formats are:

- Char (for characters)
- Varchar (for variable character length)
- Numeric (Integer, time)

## 2.1 Types of SQL Statements

In the SQL world, we can ultimately distinguish between the data definition and manipulation language:

- *Data Definition Language (DDL)*:
  - Commands from the DDL are used to define, change, or drop tables (database)
  - SQL examples: Create, Alter, Truncate, Drop
- *Data Manipulation Language (DML)*:
  - DML is used to read and modify data in tables
  - Those operations are sometimes named as CRUD operations and we learned them in the last blog: *Create, Read, Update, and Delete* rows in a table
  - SQL examples: INSERT, SELECT, UPDATE, DELETE

Now, let's put some of these concepts into practice:

## 2.2 CREATE

- You can create new tables with **CREATE TABLE** command. It works in three steps. You have to provide a name for your table, each column needs a name, and you have to specify which kind of information will be stored (e.g. numerical values, characters) in the column
- The following command creates a toy table with for petsales with five variables:

```
CREATE TABLE PETALE (
  ID INTEGER NOT NULL,
  PET CHAR(20),
  SALEPRICE DECIMAL(6,2),
  PROFIT DECIMAL(6,2),
  SALEDATE DATE
);
```

As the output illustrates, we can add options to create the table:

- The ID variable is an integer that does not accept zero, or in other words: **NOT NULL**
- The column PET is generated to store to character string
- The column SALEDATE stores dates
- And we could also set a primary key with the clause: **PRIMARY KEY**
- A second example

```
CREATE TABLE PET (
    ID INTEGER NOT NULL,
    ANIMAL VARCHAR(20),
    QUANTITY INTEGER
);
```

- So far, both tables do not contain any values. With `INSERT INTO`, we fill the table with corresponding values:

```
INSERT INTO PETALE VALUES
    (1, 'Cat', 450.09, 100.47, '2018-05-29'),
    (2, 'Dog', 666.66, 150.76, '2018-06-01'),
    (3, 'Parrot', 50.00, 8.9, '2018-06-04'),
    (4, 'Hamster', 60.60, 12, '2018-06-11'),
    (5, 'Goldfish', 48.48, 3.5, '2018-06-14');
```

- And for the second table:

```
INSERT INTO PET VALUES
    (1, 'Cat', 3),
    (2, 'Dog', 4),
    (3, 'Hamster', 2);
```

As we learned in the last session, we can use `SELECT` to check whether it worked:

```
SELECT * FROM PET;
```

Table 2.1: 3 records

ID	ANIMAL	QUANTITY
1	Cat	3
2	Dog	4
3	Hamster	2

## 2.3 ALTER TABLE

- We use the `ALTER TABLE` statement to add, delete, or modify columns. For instance:
- `ADD COLUMN`, `DROP COLUMN`; `ALTER COLUMN`; `RENAME COLUMN`

First: `ADD COLUMN`



```
ALTER TABLE PETALE
ADD COLUMN QUANTITY INTEGER;
```

```
SELECT * FROM PETALE;
```

Table 2.2: 5 records

ID	PET	SALEPRICE	PROFIT	SALEDATE	QUANTITY
1	Cat	450.09	100.47	2018-05-29	NA
2	Dog	666.66	150.76	2018-06-01	NA
3	Parrot	50.00	8.90	2018-06-04	NA
4	Hamster	60.60	12.00	2018-06-11	NA
5	Goldfish	48.48	3.50	2018-06-14	NA

- Again, fill in your values

```
UPDATE PETALE SET QUANTITY = 9 WHERE ID = 1;
```

```
UPDATE PETALE SET QUANTITY = 24 WHERE ID = 5;
```

- Check whether it worked

```
SELECT * FROM PETALE;
```

Table 2.3: 5 records

ID	PET	SALEPRICE	PROFIT	SALEDATE	QUANTITY
1	Cat	450.09	100.47	2018-05-29	9
2	Dog	666.66	150.76	2018-06-01	NA
3	Parrot	50.00	8.90	2018-06-04	NA
4	Hamster	60.60	12.00	2018-06-11	NA
5	Goldfish	48.48	3.50	2018-06-14	24

Second: DROP COLUMN

```
ALTER TABLE PETALE
DROP COLUMN PROFIT;
```

```
SELECT * FROM PETALE;
```

Table 2.4: 5 records

ID	PET	SALEPRICE	SALEDATE	QUANTITY
1	Cat	450.09	2018-05-29	9
2	Dog	666.66	2018-06-01	NA
3	Parrot	50.00	2018-06-04	NA
4	Hamster	60.60	2018-06-11	NA
5	Goldfish	48.48	2018-06-14	24

Third: ALTER COLUMN

- We can change the data type, for instance, to increase the length of a character variable to VARCHAR(20) with ALTER COLUMN

```
ALTER TABLE PETALE  
ALTER COLUMN PET SET DATA TYPE VARCHAR(20);
```

```
SELECT * FROM PETALE;
```

Table 2.5: 5 records

ID	PET	SALEPRICE	SALEDATE	QUANTITY
1	Cat	450.09	2018-05-29	9
2	Dog	666.66	2018-06-01	NA
3	Parrot	50.00	2018-06-04	NA
4	Hamster	60.60	2018-06-11	NA
5	Goldfish	48.48	2018-06-14	24

Forth: RENAME COLUMN

- Use RENAME COLUMN to change *to* a new name:

```
ALTER TABLE PETALE  
RENAME COLUMN PET TO ANIMAL;
```

```
SELECT * FROM PETALE;
```

Table 2.6: 5 records

ID	ANIMAL	SALEPRICE	SALEDATE	QUANTITY
1	Cat	450.09	2018-05-29	9
2	Dog	666.66	2018-06-01	NA
3	Parrot	50.00	2018-06-04	NA
4	Hamster	60.60	2018-06-11	NA
5	Goldfish	48.48	2018-06-14	24

## 2.4 Truncate

- The TRUNCATE statement will remove all(!) rows from an existing table, just like the one we created in the beginning, however, it does not delete the table itself.

```
TRUNCATE TABLE PET IMMEDIATE;
```

Caution: DROP TABLE tablename; drops the entire table!

```
DROP TABLE PETALE;
```

## 3 Calculations

Do you really want to “learn SQL the badass way”? I outlined the scope of the blog in the first blog entry where I explained why I set the scope to R users and people with some programming languages. You may not find what you are looking for if you are not familiar with R or if you do not have any other experience to work with data, because I focus on the data management part in this blog. Thus, I hope you find some helpful resources to learn SQL, but I focus on data wrangling aspects, without explaining basic concepts to handle data.

In this session we learn how to use string patterns and ranges to search data. We will learn how to sort and group data to display result. Moreover, we practice composing nested queries and execute select statements to access data from multiple tables. For this reason I created a simple table that contains some attributes (y, z, id) about countries:

```
SELECT * FROM df ;
```

Table 3.1: 4 records

country	y	z	id
Germany	2	3.6	1
Austria	1	8.5	2
Brazil	4	2.5	3
Brazil	3	3.5	3

### 3.1 String values, ranges and set of values

We can use strings and the WHERE clause to search for specific observations. For instance, `WHERE countryname LIKE 'A%'` means that we search for country name column that start with the corresponding string. And we can use % as a wildcard character:

```
SELECT * FROM df WHERE country LIKE 'A%';
```

Table 3.2: 1 records

country	y	z	id
Austria	1	8.5	2

Use a range to select entries that depending on some criteria ( $>$  and  $<$ ). In SQL, we specify **WHERE values are between 100 and 200**. Keep in mind that values are inclusive within the range. For instance, we can use the mtcars dataset and restrict the table with cars that have a horsepower (hp) between 100 and 200, we can even use an **AND** to restrict to cars with a manual transmission ( $AM = 1$ )

```
select * from mtcars
where (hp BETWEEN 100 and 200) and AM = 1 ;
```

Table 3.3: 5 records

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

Another option gives us the **IN** operator. We can select columns by providing a list and the **IN** operator. As the next example shows, we select only those observations that match the provided list of the **IN** operator:

```
SELECT * FROM df WHERE country IN ('Brazil');
```

Table 3.4: 2 records

country	y	z	id
Brazil	4	2.5	3
Brazil	3	3.5	3

### 3.1.1 Sorting Result Sets

Sometimes we need to sort the entries alphabetically and we can do that with the **ORDER BY** clause. For instance, **ORDER BY country**:

```
SELECT * FROM df ORDER BY country;
```

Table 3.5: 4 records

country	y	z	id
Austria	1	8.5	2
Brazil	4	2.5	3
Brazil	3	3.5	3
Germany	2	3.6	1

- By default, the entries are ordered in an ascending order, but we can sort in a descending with DESC option as well:

```
SELECT * FROM df ORDER BY country DESC;
```

Table 3.6: 4 records

country	y	z	id
Germany	2	3.6	1
Brazil	4	2.5	3
Brazil	3	3.5	3
Austria	1	8.5	2

Sometimes we have several observations per unit or any kind of structural order, which is why we may want to order a specific variable. We can sort the data by providing the number of the column we want to use a sort. As the next example shows, we can use y (or 2) to sort the data:

```
SELECT * FROM df ORDER BY y;
```

Table 3.7: 4 records

country	y	z	id
Austria	1	8.5	2
Germany	2	3.6	1
Brazil	3	3.5	3
Brazil	4	2.5	3

### 3.1.2 Grouping result sets

To work with data, we have to get rid of duplicates and often it is much more easier if we restrict result set (data frame). To exclude duplicates we can use the `distinct()` command, which returns only distinct countries in our example:

```
SELECT distinct(country) FROM df ;
```

Table 3.8: 3 records

country
Germany
Austria
Brazil

In a similar fashion, maybe we have to clarify how many observations do we have per group? Or in our case, how many entries come from the same country and how often appears each level? In such a case we can count the county column and use the `group by` clause:

```
SELECT country, count (country) from df group by country;
```

Table 3.9: 3 records

country	count (country)
Austria	1
Brazil	2
Germany	1

As the last output showed, the count functions literally counts, but SQL does not give it a name, it simply displays what it does. We can change this ugly behaviour by providing a variable name and tell SQL how the column should be listed.

```
SELECT country, count (country) AS Count_Variable from df group by country;
```

Table 3.10: 3 records

country	Count_Variable
Austria	1

country	Count_Variable
Brazil	2
Germany	1

Certainly, counting is not the only function. We can estimate the mean average with `AVG()`. And now the average, little SQL monkey!

```
SELECT country, AVG(z) as Mean from df group by country;
```

Table 3.11: 3 records

country	Mean
Austria	8.5
Brazil	3.0
Germany	3.6

We can set a further conditions with a grouped by clause and add the `HAVING` option. As the next output shows, the `group by country HAVING count (country) > 1` returns only countries with more than one observation counted:

```
SELECT country, count (country) AS Count from df group by country having count (country) >
```

Table 3.12: 1 records

country	Count
Brazil	2

Let us try to remember that the `WHERE` clause is for entire result set; while `HAVING` works only for the `GROUPED BY` clause.

Congratulations! At this point you are able to:

- Use the `WHERE` clause to refine your query results
- Use the wildcard character (%) as a substitute for unknown characters in a pattern
- Use `BETWEEN ... AND` to specify a range of numbers
- We can sort query results into ascending or descending order, by using the `ORDER BY` clause
- And we can group query results by using the `GROUP BY` clause.



## 3.2 Built-in Database Functions

We saw in the last section that we can aggregate (count, avg) data and use column functions. Most of the basic statistics functions (sum, avg, min, max) are available and we can specify further conditions, for instance, if we want to summarize groups:

```
select sum(mpg) as sum_mpg from mtcars where hp > 100
```

Table 3.13: 1 records

sum_mpg
401.4

Or we may use the scalar function and round to the nearest integer:

```
select round(drat, 1) as round_drat from mtcars
```

Table 3.14: Displaying records 1 - 10

round_drat
3.9
3.9
3.9
3.1
3.2
2.8
3.2
3.7
3.9
3.9

In SQL, there is a class of scalar functions. For instance, we can calculate the length of a string:

```
select length(country) from df
```

Table 3.15: 4 records

length(country)
7
7
6
6

Depending the SQL database you use, in db2 you can use the upper (UCASE) and lower case (LCASE) function for strings.

```
select upper(country) from df
```

Table 3.16: 4 records

upper(country)
GERMANY
AUSTRIA
BRAZIL
BRAZIL

In case of Oracle the functions are called lower and upper.

### 3.2.1 Date and Time Built-in Functions

Talking about SQL databases, there are three different possibilities to work with date and time DB2.

- Date: *YYYYMMDD* (Year/Month/Day) - Time: *HHMMSS* (Hours/Min/Sec) - Timestamp: *YYYYMMDDHHMMSSZZZZZZ* (Date/Time/Microseconds)

Depending on what you are up to do, there are functions to extract the day, month, day of month, day of week, day of year, week, hour, minute, and second. You can also extract the `current_date` and the `current_time`. Unfortunately, this does not work in Oracle the same way as in DB2, but to give you an example how to extract the day:

```
select day(date) from df where country = 'Germany'
```

### 3.2.2 Sub-Queries and Nested Selects

Sub-queries or sub selects are like regular queries but placed within parentheses and nested inside another query. This allows you to form more powerful queries than would have been otherwise possible. An example:

```
select avg(mpg) from mtcars
```

Table 3.17: 1 records

avg(mpg)
20.09062

Let's say we want to select only the observations with higher values than the average of mpg:

```
select * from mtcars where mpg > avg(mpg)
```

This would produce the following error: `misuse of aggregate function avg() Failed to execute SQL chunk`. One of the limitations of built in aggregate functions, like `avg()`, is that they cannot be evaluated in the `WHERE` clause always. Thus, we have to use sub-queries.

```
select * from mtcars where mpg >
(select avg(mpg) from mtcars);
```

Table 3.18: Displaying records 1 - 10

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1

Column expressions help to set sub queries as a list of columns. Say we select variable `Z`:

```
select country, z from df
```

Table 3.19: 4 records

country	z
Germany	3.6
Austria	8.5
Brazil	2.5
Brazil	3.5

And in the next step we add the average of all countries:

```
select country, z, avg(z) as avg_Z from df
```

Table 3.20: 1 records

country	z	avg_Z
Germany	3.6	4.525

This is obviously wrong. We cannot calculate on micro and macro level the same time, but we could use a sub-query (also called table expressions) to achieve it:

```
select country, z, (select avg(z) from df group by country) as avg_Z from df
```

Table 3.21: 4 records

country	z	avg_Z
Germany	3.6	8.5
Austria	8.5	8.5
Brazil	2.5	8.5
Brazil	3.5	8.5

Sub-queries can also be applied in the from clause. They are called derived tables or table expressions, because the outer query uses the results of the sub-query as a data source

```
select * from (select hp, vs from mtcars);
```

Table 3.22: Displaying records 1 - 10

hp	vs
110	0
110	0
93	1
110	1
175	0
105	1
245	0
62	1
95	1
123	1

### 3.2.3 Working with Multiple Tables

There are several ways to access multiple tables in the same query. Namely, using sub-queries, implicit join, and join operators, such as `INNER JOIN` and `OUTER JOIN`. For instance:

```
select * from df2;
```

Table 3.23: 2 records

country	valid	id
Germany	1	1
Austria	0	2

Let's say we want only observations from `df` that are listed in `df2`. In such a situation we can use a sub-queries:

```
select * from df
  where country in
    (select country from df2)
```

Table 3.24: 2 records

country	y	z	id
Germany	2	3.6	1

country	y	z	id
Austria	1	8.5	2

Of course, you could add also information of the second table and include only countries with a certain value:

```
select * from df
where country in
(select country from df2 where valid = 1)
```

Table 3.25: 1 records

country	y	z	id
Germany	2	3.6	1

Implicit joins implies that we can access multiple tables by specifying them in the **FROM** clause of the query. This leads to a **CROSS JOIN** (also known as Cartesian Join).

```
select * from df, df2
```

Table 3.26: 8 records

country	y	z	id	country	valid	id
Germany	2	3.6	1	Germany	1	1
Germany	2	3.6	1	Austria	0	2
Austria	1	8.5	2	Germany	1	1
Austria	1	8.5	2	Austria	0	2
Brazil	4	2.5	3	Germany	1	1
Brazil	4	2.5	3	Austria	0	2
Brazil	3	3.5	3	Germany	1	1
Brazil	3	3.5	3	Austria	0	2

In DBL2 we can use the where clause to match data (see code); in Oracle there are other matching operators

```
select * from df, df2 where df.id = df2.id;
```

In case of long names, we can use shorter aliases for table names (or use column names with aliases in the **SELECT** clause):

```
select * from df A, df2 B where A.id = B.id;
```

### 3.3 Summary

- Most databases come with built-in functions that you can use in SQL statements to perform operations on data within the database itself.
- When you work with large datasets, you may save time by using built-in functions rather than first retrieving the data into your application and then executing functions on the retrieved data.
- Use sub-queries to form more powerful queries.
- A sub-select expression helps to evaluate some built-in aggregate functions like the average function.
- Derived tables or table expressions are sub-queries where the outer query uses the results of the sub-query as a data source.

## 4 Cheats

This is the cheat sheet section of this book. It only shows example code for a copy and paste approach for each chapter.

### 4.1 Introduction

#### 4.1.1 Select

SELECT statement is used to fetch data from a database.

```
SELECT * FROM mtcars;
```

Table 4.1: Displaying records 1 - 10

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4

#### 4.1.2 Where

WHERE clause is used to extract only those records that fulfill a condition.

```
SELECT * FROM mtcars WHERE am = 0 LIMIT 5;
```



Table 4.2: 5 records

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2

### 4.1.3 Count

COUNT takes the name of a column as argument and counts the number of rows when the column is not NULL.

```
SELECT COUNT(*) FROM mtcars;
```

Table 4.3: 1 records

COUNT(*)
32

### 4.1.4 DISTINCT

Get unique values in specified columns.

```
SELECT DISTINCT Species FROM iris
```

Table 4.4: 3 records

Species
setosa
versicolor
virginica

### 4.1.5 Insert Values

Insert new rows in the table.

```
library(tidyverse)

df <- tribble(
  ~x, ~y, ~z,
  "a", 2, 3.6,
  "b", 1, 8.5
)
df
```

```
# A tibble: 2 x 3
  x     y     z
<chr> <dbl> <dbl>
1 a       2   3.6
2 b       1   8.5
```

Now, we can insert new values into `df`.

```
INSERT INTO df (x, y, z) VALUES('c', 3, 1);
SELECT * FROM df;
```

#### 4.1.6 Updates

Update the rows in the table.

```
UPDATE df SET z = 77 WHERE x = 'b';
```

#### 4.1.7 Delete

Remove rows from the table which are specified in the WHERE condition.

```
DELETE from df WHERE x = 'b';
```

## 4.2 Data management

### 4.2.1 Create table

Each column in the table is specified with its name, data type and an optional keyword which could be PRIMARY KEY, NOT NULL, etc.,

```
CREATE TABLE PETALE (
  ID INTEGER NOT NULL,
  PET CHAR(20),
  SALEPRICE DECIMAL(6,2),
  PROFIT DECIMAL(6,2),
  SALEDATE DATE
);
```

INSERT INTO fills the table:

```
INSERT INTO PETALE VALUES
(1, 'Cat', 450.09, 100.47, '2018-05-29'),
(2, 'Dog', 666.66, 150.76, '2018-06-01'),
(3, 'Parrot', 50.00, 8.9, '2018-06-04'),
(4, 'Hamster', 60.60, 12, '2018-06-11'),
(5, 'Goldfish', 48.48, 3.5, '2018-06-14');
```

## 4.2.2 Alter table

ADD COLUMN:

```
ALTER TABLE PETALE
ADD COLUMN QUANTITY INTEGER;
```

Fill in values:

```
UPDATE PETALE SET QUANTITY = 9 WHERE ID = 1;
```

Check whether it worked

```
SELECT * FROM PETALE;
```

Table 4.5: 5 records

ID	PET	SALEPRICE	PROFIT	SALEDATE	QUANTITY
1	Cat	450.09	100.47	2018-05-29	9
2	Dog	666.66	150.76	2018-06-01	NA
3	Parrot	50.00	8.90	2018-06-04	NA
4	Hamster	60.60	12.00	2018-06-11	NA
5	Goldfish	48.48	3.50	2018-06-14	NA

DROP COLUMN:

```
ALTER TABLE PETALE  
DROP COLUMN PROFIT;
```

```
SELECT * FROM PETALE;
```

Table 4.6: 5 records

ID	PET	SALEPRICE	SALEDATE	QUANTITY
1	Cat	450.09	2018-05-29	9
2	Dog	666.66	2018-06-01	NA
3	Parrot	50.00	2018-06-04	NA
4	Hamster	60.60	2018-06-11	NA
5	Goldfish	48.48	2018-06-14	NA

ALTER COLUMN:

```
ALTER TABLE PETALE  
ALTER COLUMN PET SET DATA TYPE VARCHAR(20);
```

```
SELECT * FROM PETALE;
```

Table 4.7: 5 records

ID	PET	SALEPRICE	SALEDATE	QUANTITY
1	Cat	450.09	2018-05-29	9
2	Dog	666.66	2018-06-01	NA
3	Parrot	50.00	2018-06-04	NA
4	Hamster	60.60	2018-06-11	NA
5	Goldfish	48.48	2018-06-14	NA

RENAME COLUMN:

```
ALTER TABLE PETALE  
RENAME COLUMN PET TO ANIMAL;
```

```
SELECT * FROM PETALE;
```

Table 4.8: 5 records

ID	ANIMAL	SALEPRICE	SALEDATE	QUANTITY
1	Cat	450.09	2018-05-29	9
2	Dog	666.66	2018-06-01	NA
3	Parrot	50.00	2018-06-04	NA
4	Hamster	60.60	2018-06-11	NA
5	Goldfish	48.48	2018-06-14	NA

### 4.2.3 Truncate

The `TRUNCATE` statement will remove all(!) rows from an existing table, just like the one we created in the beginning, however, it does not delete the table itself.

```
TRUNCATE TABLE PET IMMEDIATE;
```

Caution: `DROP TABLE tablename;` drops the entire table!

```
DROP TABLE PETALE;
```

## 4.3 Calculations

For this subsection I created a simple table that contains attributes about countries:

```
SELECT * FROM df ;
```

Table 4.9: 4 records

country	y	z	id
Germany	2	3.6	1
Austria	1	8.5	2
Brazil	4	2.5	3
Brazil	3	3.5	3

### 4.3.1 Like

The `LIKE` operator is used in a `WHERE` clause to search for a specified pattern in a column.

There are two wildcards often used in conjunction with the LIKE operator which are percent sign(%) and underscore sign (\_).

```
SELECT * FROM df WHERE country LIKE 'A%';
```

Table 4.10: 1 records

country	y	z	id
Austria	1	8.5	2

### 4.3.2 Between

Select data within a range:

```
select * from mtcars  
where (hp BETWEEN 100 and 200) and AM = 1 ;
```

Table 4.11: 5 records

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

### 4.3.3 In

Use the IN operator to select observations that match the provided list of the IN operator:

```
SELECT * FROM df WHERE country IN ('Brazil');
```

Table 4.12: 2 records

country	y	z	id
Brazil	4	2.5	3
Brazil	3	3.5	3

### 4.3.4 Order by

ORDER BY is used to sort the result-set in ascending or descending order. The default is ascending.

```
SELECT * FROM df ORDER BY country;
```

Table 4.13: 4 records

country	y	z	id
Austria	1	8.5	2
Brazil	4	2.5	3
Brazil	3	3.5	3
Germany	2	3.6	1

### 4.3.5 Distinct

Exclude duplicates with DISTINCT:

```
SELECT distinct(country) FROM df ;
```

Table 4.14: 3 records

country
Germany
Austria
Brazil

### 4.3.6 Group by

GROUP BY is used in collaboration with SELECT to arrange identical data into groups.

```
SELECT country, count (country) AS Var_Name from df group by country;
```

Table 4.15: 3 records

country	Var_Name
Austria	1
Brazil	2
Germany	1

### 4.3.7 Average

```
SELECT country, AVG(z) as Mean from df group by country;
```

Table 4.16: 3 records

country	Mean
Austria	8.5
Brazil	3.0
Germany	3.6

### 4.3.8 Having

The WHERE clause is for entire result set; while HAVING works only for the GROUPED BY clause.

```
SELECT country, count (country) AS Count from df group by country having count (country) >
```

Table 4.17: 1 records

country	Count
Brazil	2

### 4.3.9 Further functions

Summarize groups:

```
select sum(mpg) as sum_mpg from mtcars where hp > 100
```



Table 4.18: 1 records

sum_mpg
401.4

Round to the nearest integer:

```
select round(drat, 1) as round_drat from mtcars
```

Table 4.19: Displaying records 1 - 10

round_drat
3.9
3.9
3.9
3.1
3.2
2.8
3.2
3.7
3.9
3.9

The length of a string:

```
select length(country) from df
```

Table 4.20: 4 records

length(country)
7
7
6
6

Depending the SQL database you use, in db2 you can use the upper (UCASE) and lower case (LCASE) function for strings.

```
select upper(country) from df
```

Table 4.21: 4 records

upper(country)
GERMANY
AUSTRIA
BRAZIL
BRAZIL

In case of Oracle the functions are called lower and upper.

#### 4.3.10 Date and Time Built-in Functions

Talking about SQL databases, there are three different possibilities to work with date and time DB2.

- Date: *YYYYMMDD* (Year/Month/Day) - Time: *HHMMSS* (Hours/Min/Sec) - Timestamp: *YYYYMMDDHHMMSSZZZZZZ* (Date/Time/Microseconds)

Depending on what you are up to do, there are functions to extract the day, month, day of month, day of week, day of year, week, hour, minute, and second. You can also extract the `current_date` and the `current_time`. Unfortunately, this does not work in Oracle the same way as in DB2, but to give you an example how to extract the day:

```
select day(date) from df where country = 'Germany'
```

#### 4.3.11 Working with Multiple Tables

There are several ways to access multiple tables in the same query. Namely, using sub-queries, implicit join, and join operators, such as `INNER JOIN` and `OUTER JOIN`. For instance:

```
select * from df2;
```

Table 4.22: 2 records

country	valid	id
Germany	1	1
Austria	0	2

Let's say we want only observations from df that are listed in df2. In such a situation we can use a sub-queries:

```
select * from df
  where country in
    (select country from df2)
```

Table 4.23: 2 records

country	y	z	id
Germany	2	3.6	1
Austria	1	8.5	2

Of course, you could add also information of the second table and include only countries with a certain value:

```
select * from df
  where country in
    (select country from df2 where valid = 1)
```

Table 4.24: 1 records

country	y	z	id
Germany	2	3.6	1

Implicit joins implies that we can access multiple tables by specifying them in the **FROM** clause of the query. This leads to a **CROSS JOIN** (also known as Cartesian Join).

```
select * from df, df2
```

Table 4.25: 8 records

country	y	z	id	country	valid	id
Germany	2	3.6	1	Germany	1	1
Germany	2	3.6	1	Austria	0	2
Austria	1	8.5	2	Germany	1	1
Austria	1	8.5	2	Austria	0	2
Brazil	4	2.5	3	Germany	1	1
Brazil	4	2.5	3	Austria	0	2

country	y	z	id	country	valid	id
Brazil	3	3.5	3	Germany	1	1
Brazil	3	3.5	3	Austria	0	2

In DBL2 we can use the where clause to match data (see code); in Oracle there are other matching operators

```
select * from df, df2 where df.id = df2.id;
```

In case of long names, we can use shorter aliases for table names (or use column names with aliases in the SELECT clause):

```
select * from df A, df2 B where A.id = B.id;
```

## References

- Grothendieck, G. 2017. “Sqlf: Manipulate r Data Frames Using SQL.” <https://CRAN.R-project.org/package=sqldf>.
- Müller, Kirill, Hadley Wickham, David A. James, and Seth Falcon. 2023. “RSQLite: SQLite Interface for r.” <https://CRAN.R-project.org/package=RSQLite>.
- Ooms, Jeroen, David James, Saikat DebRoy, Hadley Wickham, and Jeffrey Horner. 2022. “RMySQL: Database Interface and ‘MySQL’ Driver for r.” <https://CRAN.R-project.org/package=RMySQL>.
- R Special Interest Group on Databases (R-SIG-DB), Hadley Wickham, and Kirill Müller. 2022. “DBI: R Database Interface.” <https://CRAN.R-project.org/package=DBI>.
- Ripley, Brian, and Michael Lapsley. 2022. “RODBC: ODBC Database Access.” <https://CRAN.R-project.org/package=RODBC>.
- Urbanek, Simon. 2022. “RJDBC: Provides Access to Databases Through the JDBC Interface.” <https://CRAN.R-project.org/package=RJDBC>.
- Wickham, Hadley, Romain François, Lionel Henry, Kirill Müller, and Davis Vaughan. 2023. “Dplyr: A Grammar of Data Manipulation.” <https://CRAN.R-project.org/package=dplyr>.